

# CS4442: Computation & Architecture

## Cloud Computing Project - Spring Semester 2024

Student: Oluwajomiloju Olajitan: 23373326

### Project Description

Upon conducting rigorous research into websites that were not only relevant, but beneficial to the task at hand, I decided to deploy an open-source website from the repository <https://github.com/Sowwyz/sowwyz-drawing>. The repository contained a drawing website without a database implemented or a dockerized image.

There are several reasons why I picked this website but the main cause was because during this project, I wanted to fully dockerize a website that I have an interest in, being art. I also chose this website as I have a history using various sites from this repository for my own personal uses. Lastly, I wanted to choose a website that had not implemented any aws features, a docker image, a running kubernetes cluster, or helm charts. With all these factors, I ended up making this decision.

There are many benefits to me choosing to containerize this website. The app is currently designed as a temporary canvas where users can draw anything, switch colours and just reset their image. With a system like this in place, I believe users should be able to interact with a database to store their images and be able to view all works of art they have made in the past. This website is also open to future expansion, as with the introduction of aws features, you may view the gallery of all users once granted permission. By using containers and managing them with Kubernetes, you can easily scale your application to meet changing demands. This improves both its availability and performance. Even if an individual container malfunctions, Kubernetes can automatically restart it, ensuring your application remains accessible to users.

### 1. Local Deployment Using Docker

The first order of operation was to create a docker file.

```
Dockerfile > CMD
1 FROM node:14
2
3 WORKDIR /app
4
5 # Copies package.json and package-lock.json files
6 COPY package.json ./
7 COPY package-lock.json ./
8
9 RUN npm install
10
11 # This copies the rest of the application code
12 COPY . .
13
14 EXPOSE 3002
15
16 CMD ["node", "server.js"]
17
```

Fig 1.1: Dockerfile Configuration

After defining my docker file, I then built my image locally and pushed it to ECR on aws by following the operations defined.

```
The push refers to repository [154260813552.dkr.ecr.us-east-1.amazonaws.com/my-repo]
An image does not exist locally with the tag: 154260813552.dkr.ecr.us-east-1.amazonaws.com/my-repo
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ docker build -t my-repo .
[+] Building 268.9s (11/11) FINISHED
=> [internal] load build definition from Dockerfile
=> [internal] load metadata for docker.io/library/node:14
=> [internal] load .dockerignore
=> [internal] load context: 2B
=> [1/6] FROM docker.io/library/node:14@sha256:a158d3b0b4e3fa813fa6c8c590b8f9a860e015ad4e59bbce5744d2f6fd8461aa
=> resolve docker.io/library/node:14@sha256:a158d3b0b4e3fa813fa6c8c590b8f9a860e015ad4e59bbce5744d2f6fd8461aa
=> sha256:2f1d7c41c74a25280bf4f00aad0a727f84518f55f6ca845bc747976c408 50.45MB / 50.45MB
=> sha256:b253a0fca708e71b60088df01de101a38a045ff7bc656e3b0fbfc7c0ccea 7.86MB / 7.86MB
=> sha256:322701bd995ccc12851a50870de0341a317011dcb9ac9df3a50c952bb131 10.00MB / 10.00MB
```

Fig 1.2: Docker Image building

```
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ docker tag my-repo:latest 154260813552.dkr.ecr.us-east-1.amazonaws.com/my-repo:latest
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ docker push 154260813552.dkr.ecr.us-east-1.amazonaws.com/my-repo:latest
The push refers to repository [154260813552.dkr.ecr.us-east-1.amazonaws.com/my-repo]
986d30193825: Pushed
4ce72a9f1c3c: Pushed
453baa92f545: Preparing
e3584b92054f: Preparing
ba7c91e66817: Preparing
0d5f5a015e5d: Preparing
3c777d951de2: Pushed
f8a91dd5fc84: Pushed
cb81227abde5: Pushed
e01a454893a9: Pushed
c45660adde37: Pushed
fe0fb3ab4a0f: Pushed
f1186e5061f2: Pushed
b2dba7477754: Pushed
latest: digest: sha256:e21beb2efc91b1a69dbf14bf326a8d1d1512e0eed88415f2160d8f267b05c106 size: 3261
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$
```

Fig 1.3: Image pushed

## 2. Set up Kubernetes Cluster

I first decided to set up my kubernetes cluster first before helm charts because it was an area I have some experience in and thought it would be more appropriate for this reason. I created my cluster on the aws website using EKS.

Clusters (1) <a href="#">Info</a>						<a href="#">Refresh</a>	<a href="#">Delete</a>	<a href="#">Add cluster</a> ▼
<input type="text" value="Filter clusters"/>						< 1 >		
	Cluster name ▲	Status ▼	Kubernetes version ▼	Support period ▼	Provider ▼			
<input type="radio"/>	<a href="#">website-cluster</a>	Active	1.29 <a href="#">Upgrade now</a>	Standard support until March 23, 2025	EKS			

Fig 2.1: EKS Cluster Setup

```

! deployment.yaml
1  replicaCount: 1
2
3  image:
4    repository: my-docker-repo/my-app
5    tag: latest
6    pullPolicy: IfNotPresent
7
8  service:
9    type: LoadBalancer
10   port: 3002
11
12  resources:
13    limits:
14      cpu: 500m
15      memory: 512Mi
16    requests:
17      cpu: 250m
18      memory: 256Mi
19
20  ingress:
21    enabled: true
22    annotations:
23      kubernetes.io/ingress.class: nginx
24    hosts:
25      - host: my-app.dev.example.com
26        paths:
27          - path: /
28            pathType: ImplementationSpecific
29

```

Fig 2.2 : deployment.yaml

```

! service.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: my-app-service
5  spec:
6    selector:
7      app: my-app
8    ports:
9      - protocol: TCP
10        port: 3002
11        targetPort: 3002
12    type: LoadBalancer
13

```

Fig 2.3: service.yaml

```

! ingress.yaml
1  apiVersion: networking.k8s.io/v1
2  kind: Ingress
3  metadata:
4    name: my-app-ingress
5    annotations:
6      kubernetes.io/ingress.class: nginx
7  spec:
8    rules:
9      - host: my-app.dev.example.com
10     http:
11       paths:
12         - path: /
13           pathType: ImplementationSpecific
14           backend:
15             service:
16               name: my-app-service
17               port:
18                 number: 3002
19

```

Fig 2.4: ingress.yaml

After deployment, I verified the application's accessibility using the `kubectl get services` command. Initially, the application wasn't reachable. To fix the issue, I used `kubectl get pods` and `kubectl get nodes` to inspect the pods and nodes in the cluster. I discovered that there weren't any running nodes available. To address this, I created a node group.

Filter Nodes by property or value					
Node name ▲	Instance type ▼	Node group ▼	Created ▼	Status ▼	
<a href="#">ip-172-31-40-142.ec2.internal</a>	t3.medium	<a href="#">web-node-group</a>	Created May 29, 2024, 22:43 (UTC+01:00)	Ready	
<a href="#">ip-172-31-8-156.ec2.internal</a>	t3.medium	<a href="#">web-node-group</a>	Created May 29, 2024, 22:45 (UTC+01:00)	Ready	
Node groups (1) Info					
<div>EditDeleteAdd node group</div>					
Group name ▲	Desired size ▼	AMI release version ▼	Launch template ▼	Status ▼	
<a href="#">web-node-group</a>	2	1.29.3-20240514 <a href="#">Update now</a>	-	Active	

Fig 2.5: Nodes running

```
[cloudshell-user@ip-10-132-42-228 ~]$ cd jjola00.github.io
[cloudshell-user@ip-10-132-42-228 jjola00.github.io]$ kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-36-134.ec2.internal       Ready    <none>   74m   v1.29.3-eks-ae9a62a
ip-172-31-8-164.ec2.internal        Ready    <none>   76m   v1.29.3-eks-ae9a62a
[cloudshell-user@ip-10-132-42-228 jjola00.github.io]$ kubectl get pods
NAME                                READY     STATUS    RESTARTS   AGE
my-app-my-chart-55f4d49657-xpnpr    1/1       Running   0           65m
[cloudshell-user@ip-10-132-42-228 jjola00.github.io]$ kubectl get services
NAME            TYPE        CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
kubernetes      ClusterIP   10.100.0.1     <none>        443/TCP     6d22h
my-app-my-chart ClusterIP   10.100.142.147 <none>        80/TCP     65m
[cloudshell-user@ip-10-132-42-228 jjola00.github.io]$ helm install my-app ./my-app-chart
-bash: helm: command not found
[cloudshell-user@ip-10-132-42-228 jjola00.github.io]$ helm -v
-bash: helm: command not found
[cloudshell-user@ip-10-132-42-228 jjola00.github.io]$ curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3
[cloudshell-user@ip-10-132-42-228 jjola00.github.io]$ cat get_helm.sh
#!/usr/bin/env bash
```

Fig 2.6: List of kubectl services

### 3. Setting up Helm

The first step I took was to install Helm.

```
[cloudshell-user@ip-10-140-116-212 ~]$ curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3
[cloudshell-user@ip-10-140-116-212 ~]$
```

Fig 3.1: First attempt of Helm Install

At first I ran into an issue where I could not install Helm due to storage, so I freed up space in aws first. After this, I reran the code to install helm and tested the version.

```
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ docker system prune
WARNING! This will remove:
- all stopped containers
- all networks not used by at least one container
- all dangling images
- unused build cache

Are you sure you want to continue? [y/N] y
Deleted build cache objects:
sqdajvsh785linri89uti78q
ljpjg1ko7x4qeyg6wf2rmgo02
w93skk7kme4mroak7dpqn2sho
zp3tywfwzkox9itkrw9w4szv

Total reclaimed space: 310.1MB
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ chmod +x get_helm.sh
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ ./get_helm.sh
Downloading https://get.helm.sh/helm-v3.15.1-linux-amd64.tar.gz
```

Fig 3.2: Re running installation once storage freed

```
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ helm create my-chart
Creating my-chart
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ ls
aws  awsliv2.zip  Dockerfile  get_helm.sh  my-chart  my-web-helm  node_modules  package.json  package-lock.json  public  server.js
[cloudshell-user@ip-10-140-116-212 jjola00.github.io]$ cd my-chart
[cloudshell-user@ip-10-140-116-212 my-chart]$ ls
charts  Chart.yaml  templates  values.yaml
[cloudshell-user@ip-10-140-116-212 my-chart]$ cd charts
[cloudshell-user@ip-10-140-116-212 charts]$ ls
[cloudshell-user@ip-10-140-116-212 charts]$ cd ..
[cloudshell-user@ip-10-140-116-212 my-chart]$ cd templates
[cloudshell-user@ip-10-140-116-212 templates]$ ls
deployment.yaml  helpers.tpl  hpa.yaml  ingress.yaml  NOTES.txt  serviceaccount.yaml  service.yaml  tests
[cloudshell-user@ip-10-140-116-212 templates]$
```

Fig 3.3: Project setup after installing helm

After setting up and testing helm, I decided to update some of my yaml files:

```

! values.yaml
1  replicaCount: 1
2
3  image:
4    repository: my-docker-repo/my-app
5    tag: latest
6    pullPolicy: IfNotPresent
7
8  service:
9    type: LoadBalancer
10   port: 3002
11
12  resources:
13    limits:
14      cpu: 500m
15      memory: 512Mi
16    requests:
17      cpu: 250m
18      memory: 256Mi
19
20  ingress:
21    enabled: true
22    annotations:
23      kubernetes.io/ingress.class: nginx
24    hosts:
25      - host: my-app.example.com
26        paths:
27          - path: /
28            pathType: ImplementationSpecific
29

```

Fig 3.4: values.yaml

```

! serviceTEMPLATE.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: my-service
5    labels:
6      app: my-app
7  spec:
8    selector:
9      app: my-app
10   ports:
11     - protocol: TCP
12       port: 80
13       targetPort: 8080
14   type: LoadBalancer
15

```

Fig 3.5: serviceTEMPLATES.yaml

```
! deploymentTEMPLATE.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: my-deployment
5    labels:
6      app: my-app
7  spec:
8    replicas: 3
9    selector:
10     matchLabels:
11       app: my-app
12    template:
13     metadata:
14       labels:
15         app: my-app
16     spec:
17     containers:
18     - name: my-container
19       image: my-docker-repo/my-app:latest
20       ports:
21       - containerPort: 8080
22
```

Fig 3.6: deploymentTEMPLATE.yaml

```

! ingressTEMPLATE.yaml
1  apiVersion: networking.k8s.io/v1
2  kind: Ingress
3  metadata:
4    name: my-ingress
5    annotations:
6      nginx.ingress.kubernetes.io/rewrite-target: /
7  spec:
8    rules:
9      - host: my-app.example.com
10      http:
11        paths:
12          - path: /
13            pathType: Prefix
14            backend:
15              service:
16                name: my-service
17                port:
18                  number: 80
19

```

Fig 3.7: *ingresTEMPLATES.yaml*

Finally, I deployed my web chart using helm

```

[cloudshell-user@ip-10-140-116-212 templates]$ helm repo add stable https://charts.helm.sh/stable
"stable" already exists with the same configuration, skipping
[cloudshell-user@ip-10-140-116-212 templates]$ helm repo update
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "stable" chart repository
Update Complete. *Happy Helming!*
[cloudshell-user@ip-10-140-116-212 templates]$ helm install my-app ./my-chart

```

Fig 3.8: *Helm web chart deployment*

## 4. AWS SERVICE 1: S3 Buckets

The first service I successfully implemented was an S3 bucket to store images.



```
const s3 = new S3Client({  
  region: "us-east-1"  
});
```

*Fig 4.1: Initialising bucket*

Once the user draws anything on the website, they have a prompt to save it. This is then saved in buckets on the AWS.



*Fig 4.2: Website prompt to Save image*

```

app.post('/save-drawing', async (req, res) => {
  const { image } = req.body;
  const imageId = uuidv4();
  const base64Data = Buffer.from(image.replace(/^data:image\/\w+;base64/, ''), 'base64');
  const type = image.split(';')[0].split('/')[1];

  const putObjectParams = {
    Bucket: S3_BUCKET,
    Key: `${imageId}.${type}`,
    Body: base64Data,
    ContentEncoding: 'base64',
    ContentType: `image/${type}`
  };

  try {
    await s3.send(new PutObjectCommand(putObjectParams));
    const dynamoParams = {
      TableName: TABLE_NAME,
      Item: {
        ImageID: { S: imageId },
        url: { S: `https://${S3_BUCKET}.s3.amazonaws.com/${imageId}.${type}` }
      }
    };
    await ddb.send(new PutItemCommand(dynamoParams));
    res.json({ success: true, imageUrl: dynamoParams.Item.url.S });
  }
});

```

Fig 4.3: Back end logic to connect buckets

	Name ▲	AWS Region ▼	IAM Access Analyzer	Creation date ▼
<input type="radio"/>	<a href="#">drawing-gallery-bucket</a>	US East (N. Virginia) us-east-1	<a href="#">View analyzer for us-east-1</a>	May 24, 2024, 00:18:46 (UTC+01:00)
<input type="radio"/>	<a href="#">webbingbucket</a>	US East (N. Virginia) us-east-1	<a href="#">View analyzer for us-east-1</a>	May 27, 2024, 23:10:03 (UTC+01:00)

Fig 4.4: Sample Buckets. Drawing-gallery-buckets stores the images



## 5. AWS SERVICE 2: DynamoDB

After storing the buckets in the AWS, I also wanted users to be able to see their images. To accomplish this, I implemented a database using dynamoDB.

```
const ddb = new DynamoDBClient({  
  region: "us-east-1"  
});
```

Fig 5.1: Initialising DynamoDB

Once the user selects the Gallery button, A list of the picture of the images, its creator and time created will be returned.

```
app.get('/gallery', async (req, res) => {  
  const params = {  
    TableName: TABLE_NAME  
  };  
  
  try {  
    const data = await ddb.send(new ScanCommand(params));  
    const images = data.Items.map(item => item.url.S);  
    res.json({ success: true, images });  
  } catch (error) {  
    console.error(error);  
    res.status(500).json({ success: false, message: error.message });  
  }  
});  
  
app.listen(PORT, () => {  
  console.log(`Server is running on http://localhost:${PORT}`);  
});
```

Fig 5.2: Backend logic

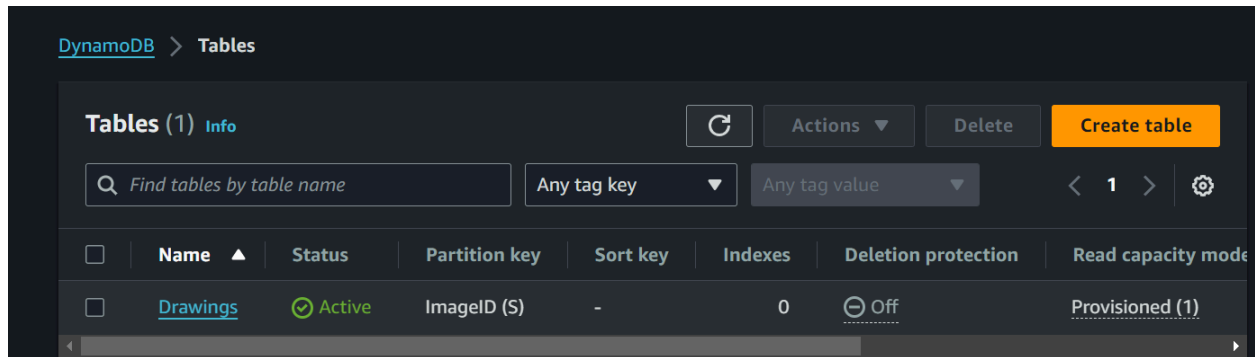


Fig 5.3: Proof of drawing table




Fig 5.4: Drawing count

## 6. Reflection

I personally found this project to be one of the most challenging of the Immersive Software Engineering syllabus so far, even after fully completing the AWS online learners course. I felt lost during various stages of the project and felt a severe lack of online resources to turn to when a roadblock is hit. However, upon completing this project I can say with utmost certainty that my software skills in relation to AWS and cloud computing has vastly improved.

One of the most important lessons I learnt during this project was the importance of going into a massive amount of work such as this with a pre-planned schedule to never veer off track and maintain tidiness and cohesiveness. The distribution of work along the 3 week period was not a linear line. The line of work and rest fluctuated due to lack of work on some days and too much work the next day. I found implementing a schedule during the latter days of the project put me at a great advantage as it allowed me to cohesively wrap up my project and tidy up any errors.

In conclusion, this project was a challenging but rewarding learning experience. It exposed knowledge gaps and emphasised the importance of continuous learning and adaptation. By analysing the challenges and my achievements, I've gained valuable insights that will propel my future projects and professional development.